

Transpiler un modèle en C : régression linéaire

December 15, 2021

1 Imports

```
[1]: import pandas as pd
import joblib
import os
from sklearn.linear_model import LinearRegression
```

2 Chargement du dataset et du modèle

```
[2]: # Chargement du dataset
df = pd.read_csv("tumors.csv")
print(df.head())

X = df[["size", "p53_concentration"]]
y = df["is_cancerous"]

# Séparation du dataset en training et testing sets
X_train = X[:-10]
X_test = X[-10:]

y_train = y[:-10]
y_test = y[-10:]

# Chargement du modèle
model = LinearRegression()

model.fit(X_train, y_train)

# Sauvegarde du modèle
joblib.dump(model, f"model.joblib")
```

	size	p53_concentration	is_cancerous
0	-0.004165	0.001785	1
1	0.012898	0.001899	1

2	0.013674	0.001193	1
3	0.008774	0.003673	0
4	0.009751	0.005571	0

[2]: ['model.joblib']

```
[3]: def produce_linear_regression_c_code():

    model = joblib.load('model.joblib')

    # Thetas
    n_thetas = len(model.coef_) + 1
    thetas = f"{model.intercept_}f,"
    for coef in model.coef_:
        thetas += str(coef) + "f,"
    thetas = thetas.strip(",")

    prediction_code = f"float thetas[{n_thetas}] = {{{thetas}}};"

    # Features
    features=""
    for i in range(X_test.shape[0]):
        to_predict = X_test.iloc[i].tolist()
        feature = "{"
        for value in to_predict:
            feature += str(value) + "f,"
        features += feature[:-2]
        features += "},\n"

    n_sample = X_test.shape[0]
    n_feature = X_test.shape[1]

    # Code
    code = f"""
#include <stdio.h>

{prediction_code}
float prediction(float *features, int n_feature)
{{
    float res = thetas[0];

    for (int i = 0; i < n_feature; ++i)
        res += features[i] * thetas[i+1];

    return res;
}}
int main()
{{
```

```

float features[{n_sample}][{n_feature}] = {{{features}}};

for (int i = 0; i < {n_sample}; ++i) {{
    printf("%f\\n", prediction(features[i], 2));
}}

return 0;
}}
"""

with open("transpiler.c", "w") as f:
    f.write(code)

```

[4]: `produce_linear_regression_c_code()`

```
!gcc transpiler.c -O3 -o transpiler
```

[5]: `print("Modèle transpilé:")`

```
!./transpiler
```

Modèle transpilé:

```

0.485140
-1.236517
0.417540
0.787183
0.768821
0.747259
-0.069009
0.386354
0.679163
0.804862

```

[6]: `print('Modèle non-transpilé:')`
`for i in model.predict(X_test):`
 `print(i)`

Modèle non-transpilé:

```

0.48513968527496
-1.2365165909081623
0.4175400010067567
0.7871834652335001
0.7688207284138135
0.7472595589533835
-0.06900933399957943
0.38635404526424
0.6791629943084592
0.8048620370086736

```

Les prédictions produites par le modèle transpilé sont bien conformes.